




Pretraining curriculum, same masking mechanism throughout

stage 1 – MAE: fixed mask rate, plain cross-entropy

stage 2 – transition: mask rate ramps linearly

stage 3 – diffusion: $t \sim U(0,1)$, reweighted loss