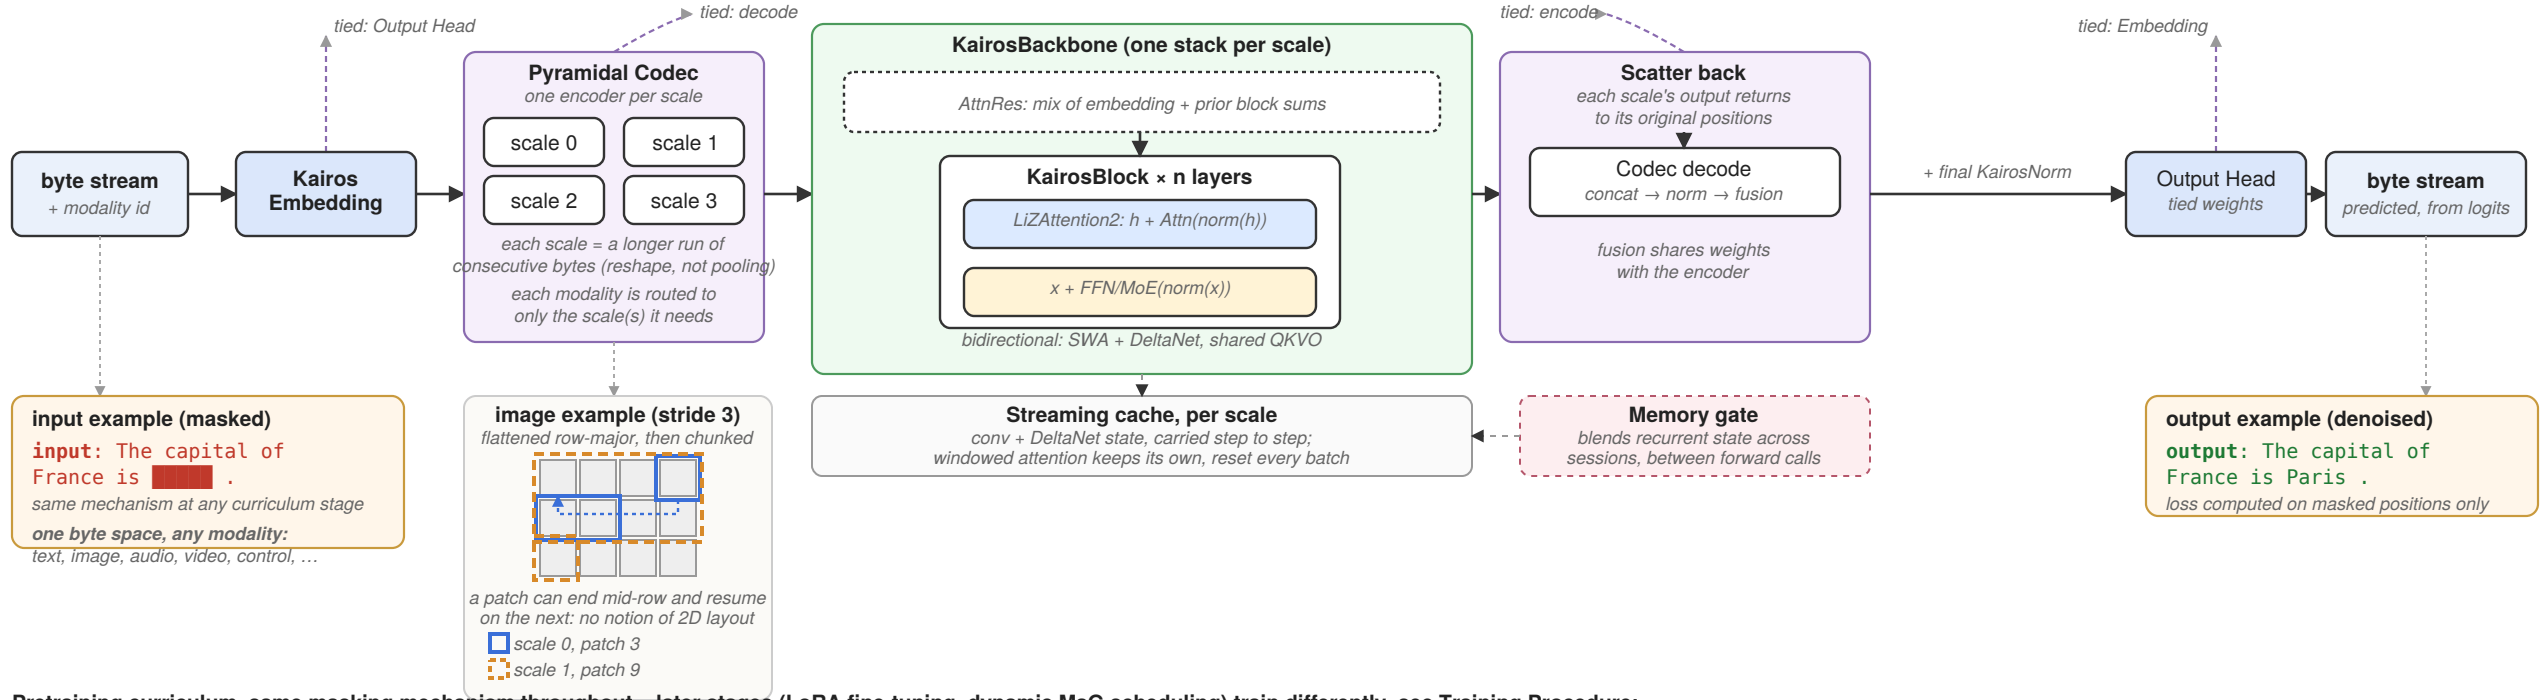



**Pretraining curriculum, same masking mechanism throughout – later stages (LoRA fine-tuning, dynamic MaG scheduling) train differently, see Training Procedure:**

stage 1 – MAE: fixed mask rate, plain cross-entropy → stage 2 – transition: mask rate ramps linearly → stage 3 – diffusion: t ~ U(0,1), reweighted loss